

Time Series Forecasting with ARIMA

Time Series Forecasting is a **crucial tool** for making informed decisions in various fields such as **weather prediction, sales analysis, business planning, and stock market analysis**. One popular method for Time Series Forecasting is the **ARIMA model**, short for **Autoregressive Integrated Moving Average**. In this article, we'll delve into the process of Time Series Forecasting using ARIMA with the **Python programming language**.

Understanding ARIMA: ARIMA is an acronym for **Autoregressive Integrated Moving Average**. It's a **powerful algorithm** specifically designed for forecasting **Time Series Data**. ARIMA models are defined by three key parameters denoted as **ARIMA(p, d, q)**, where:

- 'p' signifies the **number of lagged observations** included in the model, capturing the **autoregressive component**.
- 'd' represents the **degree of differencing** needed to make the time series **stationary**. A value of **0** indicates stationary data, while **1** suggests **seasonal data**.
- 'q' indicates the **size of the moving average window**, capturing the **moving average component** of ARIMA.

By grasping these parameters and understanding how they interact, we can effectively employ ARIMA models to generate accurate forecasts for **time-dependent datasets**. Let's explore this further by implementing ARIMA in Python for Time Series Forecasting.

To embark on Time Series Forecasting with ARIMA, our initial step involves retrieving historical data on Google's stock prices using the Yahoo Finance API. This API offers a comprehensive source of financial data, enabling us to collect a detailed dataset spanning a defined timeframe. This dataset forms the cornerstone of our analysis, facilitating insights into Google's stock price dynamics over time.

#Below is a Python code snippet demonstrating how you can retrieve the latest stock price data using the Yahoo Finance API

```
import pandas as pd
import yfinance as yf
import datetime
from datetime import date, timedelta
today = date.today()

d1 = today.strftime("%Y-%m-%d")
end_date = d1
d2 = date.today() - timedelta(days=360)
d2 = d2.strftime("%Y-%m-%d")
start_date = d2

data = yf.download('AAPL',
                   start=start_date,
                   end=end_date,
                   progress=False)

print(data.head())
```

	Open	High	Low	Close	Adj Close	\
Date						
2023-03-29	159.369995	161.050003	159.350006	160.770004	159.916428	
2023-03-30	161.529999	162.470001	161.270004	162.360001	161.497971	
2023-03-31	162.440002	165.000000	161.910004	164.899994	164.024475	
2023-04-03	164.270004	166.289993	164.220001	166.169998	165.287750	
2023-04-04	166.600006	166.839996	165.110001	165.630005	164.750626	
Volume						
Date						
2023-03-29	51305700					
2023-03-30	49501700					
2023-03-31	68749800					
2023-04-03	56976200					
2023-04-04	46278300					



The provided code retrieves stock price data spanning from the current date back to the previous 360 days. However, in this dataset, the Date information is not presented as a separate column; instead, it serves as the index. To facilitate the utilization of this data for various data science tasks, it's essential to transform this index into a regular column. Below are the steps demonstrating how to accomplish this:

```
data["Date"] = data.index
data = data[["Date", "Open", "High",
            "Low", "Close", "Adj Close", "Volume"]]
data.reset_index(drop=True, inplace=True)
print(data.head())
```

	Date	Open	High	Low	Close	Adj Close	\
0	2023-03-29	159.369995	161.050003	159.350006	160.770004	159.916428	
1	2023-03-30	161.529999	162.470001	161.270004	162.360001	161.497971	
2	2023-03-31	162.440002	165.000000	161.910004	164.899994	164.024475	
3	2023-04-03	164.270004	166.289993	164.220001	166.169998	165.287750	
4	2023-04-04	166.600006	166.839996	165.110001	165.630005	164.750626	

	Volume
0	51305700
1	49501700
2	68749800
3	56976200
4	46278300

The resultant dataset mirrors the format typically obtained from Yahoo Finance, providing comprehensive stock price data accessible through Python. This method effectively retrieves stock price information, aligning with the structure commonly encountered in datasets acquired from Yahoo Finance.

```
import pandas as pd
import yfinance as yf
import datetime
from datetime import date, timedelta
today = date.today()

d1 = today.strftime("%Y-%m-%d")
end_date = d1
d2 = date.today() - timedelta(days=365)
d2 = d2.strftime("%Y-%m-%d")
start_date = d2

data = yf.download('GOOG',
                  start=start_date,
                  end=end_date,
                  progress=False)

data["Date"] = data.index
data = data[["Date", "Open", "High", "Low", "Close", "Adj Close", "Volume"]]
data.reset_index(drop=True, inplace=True)
print(data.tail())
```

	Date	Open	High	Low	Close	Adj Close	\
246	2024-03-18	149.369995	152.929993	148.139999	148.479996	148.479996	
247	2024-03-19	148.979996	149.619995	147.009995	147.919998	147.919998	
248	2024-03-20	148.789993	149.759995	147.664993	149.679993	149.679993	
249	2024-03-21	150.320007	151.304993	148.009995	148.740005	148.740005	
250	2024-03-22	150.190002	152.550003	150.089996	151.770004	151.770004	

	Volume
246	47676700
247	17748400
248	17730000
249	19843900
250	19207179



We now proceed to select and isolate the crucial columns, 'Date' and 'Close' prices, from the dataset. This focused selection is paramount for our subsequent analysis, ensuring that we work exclusively with the essential data elements required for Time Series Forecasting and other data science tasks. By prioritizing these columns, we optimize the efficiency of our analysis, maintaining a sharp focus on the most pertinent information. This streamlined approach enhances the clarity and effectiveness of our data processing and modeling efforts, setting a solid foundation for robust insights and informed decision-making.

```
data = data[["Date", "Close"]]  
print(data.head())
```

	Date	Close
0	2023-03-24	106.059998
1	2023-03-27	103.059998
2	2023-03-28	101.360001
3	2023-03-29	101.900002
4	2023-03-30	101.320000

```
import matplotlib.pyplot as plt  
plt.style.use('fivethirtyeight')  
plt.figure(figsize=(15, 10))  
plt.plot(data["Date"], data["Close"])
```

[<matplotlib.lines.Line2D at 0x7e458339bd00>]



✓ Using ARIMA for Time Series Forecasting

Before applying the ARIMA model for Time Series Forecasting, it's essential to determine whether our dataset exhibits stationarity or seasonality. The visualization of the closing stock prices graph above indicates that our dataset lacks stationarity. To conduct a thorough assessment of stationarity and seasonality within our dataset, we can employ the seasonal decomposition method. This technique decomposes the time series data into distinct components, namely trend, seasonality, and residuals, providing valuable insights into the underlying patterns of the time series data.


- **Assessment of Data:** Before proceeding with modeling, it's crucial to evaluate the stationarity or seasonality of the dataset.
- **Utilizing Seasonal Decomposition:** The seasonal decomposition method is employed to dissect the time series data, revealing its inherent components: trend, seasonality, and residuals.
- **Insight Generation:** By decomposing the data, we gain a deeper understanding of its underlying patterns, enabling more informed modeling decisions for Time Series Forecasting.

```
from statsmodels.tsa.seasonal import seasonal_decompose
# Determine an appropriate period for seasonality based on the frequency of the data
# For daily data, we can try periods of 7 for weekly seasonality, or 30 for monthly seasonality
result = seasonal_decompose(data["Close"], model='multiplicative', period=30)

# Plot the decomposed components
fig = plt.figure()
fig = result.plot()
fig.set_size_inches(15, 10)
```



<Figure size 640x480 with 0 Axes>



Our dataset's seasonality prompts the use of the Seasonal ARIMA (SARIMA) model for Time Series Forecasting. However, we'll initially explore the ARIMA model to grasp both approaches comprehensively.

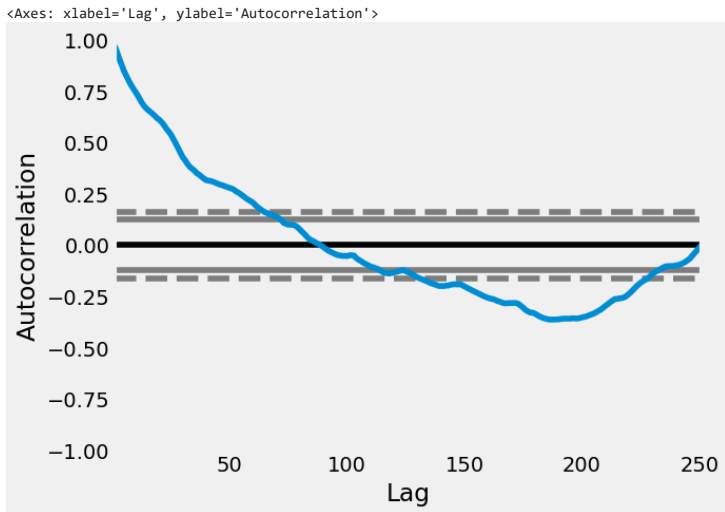
Parameter Determination:

- **p Determination:** We identify 'p' by examining the autocorrelation of the 'Close' column.
- **q Determination:** The value of 'q' is derived from the partial autocorrelation plot.
- **d Selection:** 'd' is set to 1 for seasonal data; 0 for stationary data.

This process ensures that we select suitable parameters for the ARIMA or SARIMA model, paving the way for effective Time Series Forecasting.

```
pd.plotting.autocorrelation_plot(data["Close"])
```

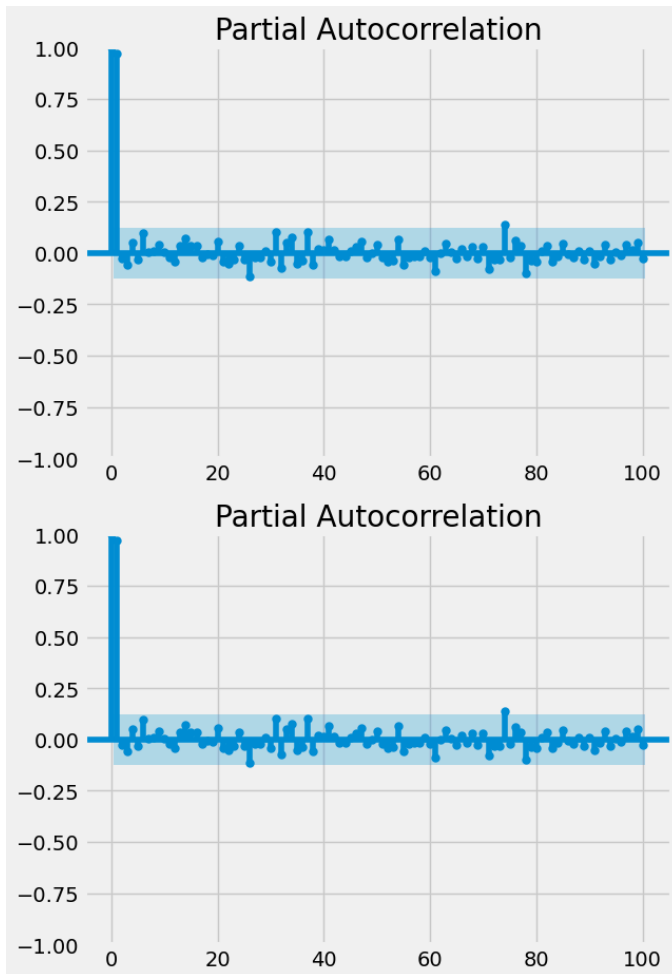




Based on the autocorrelation plot provided, we observe the curve descending after the 5th line of the first boundary. This delineates our choice for the p-value, which is determined as 5. Now, we proceed to ascertain the value of q (moving average):

```
from statsmodels.graphics.tsaplots import plot_pacf
plot_pacf(data["Close"], lags = 100)
```





From the partial autocorrelation plot presented, we discern that only 1 point extend significantly beyond the others. This characteristic informs our determination of the q value, which we identify as 2. With both the p and q values established, let's proceed to construct an ARIMA model.

```
p, d, q = 5, 1, 2
```

```
# Import ARIMA from the new module
from statsmodels.tsa.arima.model import ARIMA
```

```
# Create ARIMA model
model = ARIMA(data["Close"], order=(p, d, q))
```

```
# Fit the model
fitted = model.fit()
```

```
# Display model summary
print(fitted.summary())
```



```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:966: UserWarning: Non-stationary
warn('Non-stationary starting autoregressive parameters')
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible
warn('Non-invertible starting MA parameters found.')
SARIMAX Results
=====
Dep. Variable:          Close      No. Observations:          251
Model:                ARIMA(5, 1, 2)      Log Likelihood          -557.557
Date:                 Sat, 23 Mar 2024      AIC                   1131.114
Time:                 00:08:42      BIC                   1159.286
Sample:               0      HQIC                   1142.452
                  - 251
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          1.2390         0.117       10.610       0.000         1.010         1.468
ar.L2         -0.8830         0.155       -5.706       0.000        -1.186        -0.580
ar.L3         -0.0964         0.134       -0.717       0.473         -0.360         0.167
ar.L4          0.1566         0.134        1.165       0.244         -0.107         0.420
ar.L5         -0.1394         0.090       -1.554       0.120         -0.315         0.036
ma.L1         -1.2306         0.096      -12.860       0.000        -1.418        -1.043
ma.L2          0.9025         0.108        8.343       0.000         0.691         1.115
sigma2         5.0591         0.270       18.759       0.000         4.531         5.588
=====
Ljung-Box (L1) (Q):                0.03      Jarque-Bera (JB):          591.55
Prob(Q):                          0.86      Prob(JB):              0.00
Heteroskedasticity (H):            1.50      Skew:                  -1.25
Prob(H) (two-sided):              0.07      Kurtosis:              10.11
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```



```
predictions = fitted.predict()
print(predictions)

0      0.000000
1     106.060008
2     103.003640
3     101.246225
4     102.103144
...
246    141.868696
247    147.644047
248    147.889796
249    148.955205
250    149.246772
Name: predicted_mean, Length: 251, dtype: float64
```

Building an ARIMA model on seasonal time series data often yields inaccurate predictions due to its inability to effectively capture seasonal patterns. To address this limitation, we turn to the Seasonal ARIMA (SARIMA) model, specifically designed to handle such data.

Here's the approach to construct a SARIMA model:

```
import statsmodels.api as sm
import warnings
model=sm.tsa.statespace.SARIMAX(data['Close'],
                                order=(p, d, q),
                                seasonal_order=(p, d, q, 12))

model=model.fit()
print(model.summary())

/usr/local/lib/python3.10/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood
warnings.warn("Maximum Likelihood optimization failed to "
SARIMAX Results
=====
Dep. Variable:          Close      No. Observations:          251
Model:                SARIMAX(5, 1, 2)x(5, 1, 2, 12)      Log Likelihood          -543.069
Date:                 Sat, 23 Mar 2024      AIC                   1116.139
Time:                 00:09:58      BIC                   1168.223
```



```
Sample:                                0      HQIC                                1137.130
              - 251
              opg
Covariance Type:
=====
              coef      std err              z      P>|z|      [0.025      0.975]
-----
ar.L1      1.2329      0.167              7.381      0.000      0.906      1.560
ar.L2     -0.7925      0.167             -4.753      0.000     -1.119     -0.466
ar.L3     -0.1118      0.141             -0.793      0.428     -0.388      0.165
ar.L4      0.1497      0.144              1.041      0.298     -0.132      0.431
ar.L5     -0.1598      0.099             -1.616      0.106     -0.354      0.034
ma.L1     -1.2460      0.164             -7.587      0.000     -1.568     -0.924
ma.L2      0.8350      0.132              6.329      0.000      0.576      1.094
ar.S.L12   -0.8961      2.441             -0.367      0.714     -5.681      3.889
ar.S.L24   -0.0736      0.376             -0.196      0.845     -0.811      0.664
ar.S.L36   -0.1441      0.165             -0.873      0.383     -0.468      0.179
ar.S.L48   -0.1626      0.473             -0.344      0.731     -1.090      0.765
ar.S.L60   -0.0174      0.221             -0.078      0.937     -0.451      0.417
ma.S.L12   -0.2038      2.602             -0.078      0.938     -5.304      4.896
ma.S.L24   -0.7642      2.268             -0.337      0.736     -5.209      3.681
sigma2      4.7634      2.222              2.144      0.032      0.409      9.118
=====
Ljung-Box (L1) (Q):              0.00      Jarque-Bera (JB):              396.06
Prob(Q):              0.99      Prob(JB):              0.00
Heteroskedasticity (H):          1.38      Skew:              -0.87
Prob(H) (two-sided):          0.16      Kurtosis:              9.08
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

Let's proceed to forecast future stock prices for the upcoming 10 days using the SARIMA model.

```
predictions = model.predict(len(data), len(data)+10)
print(predictions)

251      151.784203
252      152.056600
253      152.043783
254      152.372362
255      151.242114
```